

Assignment-4

localhost8891/notebooks/Machine_learning_Assignment-4.ipynb#

jupyter Machine_learning_Assignment-4 Last Checkpoint: 9 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

```
In [34]: #importing the required libraries to work with Tabular data and also to implement algorithms

import warnings
import numpy as np
import pandas as pd
import seaborn as sns
from sklearn import preprocessing
import matplotlib.pyplot as plt
from scipy.stats.stats import pearsonr
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, recall_score, precision_score, classification_report, confusion_matrix
warnings.filterwarnings("ignore")
```

Question: 1

1. Read the provided CSV file 'data.csv'. <https://drive.google.com/drive/folders/1h8C3mlsso-R-sIOLsvoYwPLzy2fJ4IOF?usp=sharing>
2. Show the basic statistical description about the data.
3. Check if the data has null values. a. Replace the null values with the mean
4. Select at least two columns and aggregate the data using: min, max, count, mean.
5. Filter the dataframe to select the rows with calories values between 500 and 1000.
6. Filter the dataframe to select the rows with calories values > 500 and pulse < 100.
7. Create a new "df_modified" dataframe that contains all the columns from df except for "Maxpulse".
8. Delete the "Maxpulse" column from the main df dataframe
9. Convert the datatype of Calories column to int datatype.
10. Using pandas create a scatter plot for the two columns (Duration and Calories).

jupyter Machine_learning_Assignment-4 Last Checkpoint: 10 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

```
In [36]: #1. Read the provided csv file 'data.csv'. https://drive.google.com/drive/folders/1h8C3mlsso-R-sIOLsvoYwPLzy2fJ4IOF?usp=sharing

df = pd.read_csv("data.csv")
df.head()
```

Out[36]:

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0

```
In [37]: #2. Show the basic statistical description about the data.

df.describe()
```

Out[37]:

	Duration	Pulse	Maxpulse	Calories
count	169.000000	169.000000	169.000000	164.000000
mean	63.846154	107.461538	134.047337	375.790244
std	42.299649	14.510259	16.450434	266.379919
min	15.000000	80.000000	100.000000	50.300000
25%	45.000000	100.000000	124.000000	250.925000
50%	60.000000	105.000000	131.000000	318.600000
75%	60.000000	111.000000	141.000000	387.600000
max	300.000000	159.000000	184.000000	1860.400000

```
In [38]: #3. Check if the data has null values
```

```
File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

In [38]: #3. Check if the data has null values.
df.isnull().any()

Out[38]: Duration    False
Pulse             False
Maxpulse          False
Calories          True
dtype: bool

In [39]: #Replace the null values with the mean
df.fillna(df.mean(), inplace=True)
df.isnull().any()

Out[39]: Duration    False
Pulse             False
Maxpulse          False
Calories          False
dtype: bool

In [40]: #4. Select at Least two columns and aggregate the data using: min, max, count, mean.
df.agg({'Maxpulse':['min','max','count','mean'],'Calories':['min','max','count','mean']})

Out[40]:
      Maxpulse  Calories
min  100.000000  50.300000
max  184.000000 1860.400000
count 169.000000 169.000000
mean  134.047337  375.790244

In [41]: #5. Filter the dataframe to select the rows with calories values between 500 and 1000.
```

```
File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

In [41]: #5. Filter the dataframe to select the rows with calories values between 500 and 1000.
df.loc[(df['Calories']>500)&(df['Calories']<1000)]

Out[41]:
   Duration  Pulse  Maxpulse  Calories
61         80    123      146      643.1
62        160    109      135      853.0
65        180     90      130      800.4
66        150    105      135      873.4
67        150    107      130      816.0
72         90    100      127      700.0
73        150     97      127      953.2
75         90     98      125      563.2
78        120    100      130      500.4
90        180    101      127      600.1
99         90     93      124      604.1
103        90     90      100      500.4
106       180     90      120      800.3
108        90     90      120      500.3

In [42]: #6. Filter the dataframe to select the rows with calories values > 500 and pulse < 100.
df.loc[(df['Calories']>500)&(df['Pulse']<100)]

Out[42]:
   Duration  Pulse  Maxpulse  Calories
65        180     90      130      800.4
```

Greenshot
Greenshot
Exported to: Document1

JupyterLab machine_learning_Assignment-4 Last Checkpoint: 11 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

In [42]: #6. Filter the dataframe to select the rows with calories values > 500 and pulse < 100.

```
df.loc[(df['Calories']>500)&(df['Pulse']<100)]
```

Out[42]:

	Duration	Pulse	Maxpulse	Calories
65	180	90	130	800.4
70	150	97	129	1115.0
73	150	97	127	953.2
75	90	98	125	563.2
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

In [43]: #7. Create a new "df_modified" dataframe that contains all the columns from df except for "Maxpulse".

```
df_modified = df[['Duration', 'Pulse', 'Calories']]
df_modified.head()
```

Out[43]:

	Duration	Pulse	Calories
0	60	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4
4	45	117	406.0

In [44]: #8. Delete the "Maxpulse" column from the main df dataframe

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

In [44]: #8. Delete the "Maxpulse" column from the main df dataframe

```
del df['Maxpulse']
```

In [45]: df.head()

Out[45]:

	Duration	Pulse	Calories
0	60	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4
4	45	117	406.0

In [46]: df.dtypes

Out[46]:

```
Duration    int64
Pulse      int64
Calories   float64
dtype: object
```

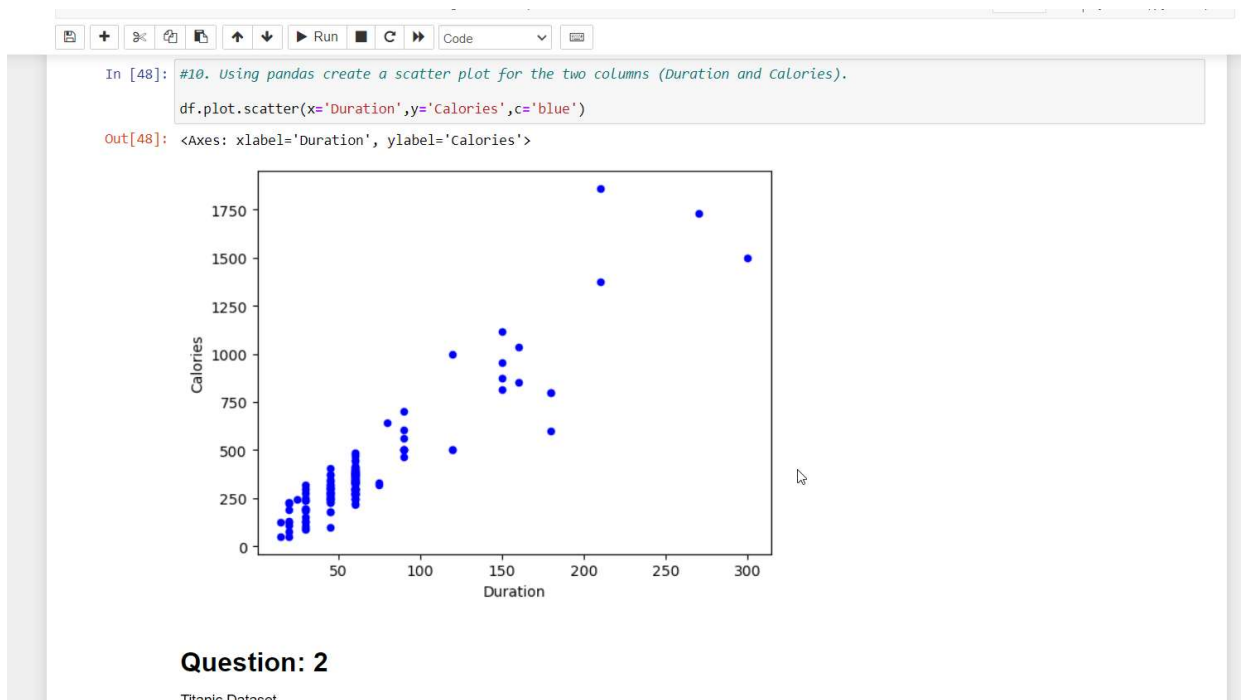
In [47]: #9. Convert the datatype of Calories column to int datatype.

```
df['Calories'] = df['Calories'].astype(np.int64)
df.dtypes
```

Out[47]:

```
Duration    int64
Pulse      int64
Calories   int64
dtype: object
```

In [48]: #10. Using pandas create a scatter plot for the two columns (Duration and Calories).



JupyterLab machine_learning_Assignment Last checkpoint: 12 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

Titanic Dataset

1. Find the correlation between 'survived' (target column) and 'sex' column for the Titanic use case in class. a. Do you think we should keep this feature?
2. Do at least two visualizations to describe or show correlations.
3. Implement Naïve Bayes method using scikit-learn library and report the accuracy

In [49]: #Loading the data file into the program

```
df=pd.read_csv("train.csv")
df.head()
```

Out[49]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cummings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

In [50]: #converted categorical data to numerical values for correlation calculation

```
label_encoder = preprocessing.LabelEncoder()
df['sex'] = label_encoder.fit_transform(df.Sex.values)

#calculation of correlation for 'Survived' and 'Sex' in data
correlation_value= df['Survived'].corr(df['sex'])

print(correlation_value)
```

-0.543351380657755

the accuracy is just 54%.so it is not necessary

Jupyter Machine_learning_Assignment-4 Last Checkpoint: 12 minutes ago (autosaved) Python 3 (ipykernel)

File Edit View Insert Cell Kernel Widgets Help Trusted

Ans: the accuracy is just 54%.so it is not necessary

```
In [51]: #print correlation matrix
matrix = df.corr()
print(matrix)
```

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare
PassengerId	1.000000	-0.005007	-0.035144	0.042939	0.036847	-0.057527	-0.001652	0.012658
Survived	-0.005007	1.000000	-0.338481	-0.543351	-0.077221	-0.035322	0.081629	0.257307
Pclass	-0.035144	-0.338481	1.000000	0.131900	-0.369226	0.083081	0.018443	-0.549500
Sex	0.042939	-0.543351	0.131900	1.000000	0.093254	-0.114631	-0.245489	-0.182333
Age	0.036847	-0.077221	-0.369226	0.093254	1.000000	-0.308247	-0.189119	0.096067
SibSp	-0.057527	-0.035322	0.083081	-0.114631	-0.308247	1.000000	0.414838	0.159651
Parch	-0.001652	0.081629	0.018443	-0.245489	-0.189119	0.414838	1.000000	0.216225
Fare	0.012658	0.257307	-0.549500	-0.182333	0.096067	0.159651	0.216225	1.000000

```
In [52]: # One way of visualizing correlation matrix in form of spread chart
df.corr().style.background_gradient(cmap="Reds")
```

Out[52]:

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare
PassengerId	1.000000	-0.005007	-0.035144	0.042939	0.036847	-0.057527	-0.001652	0.012658
Survived	-0.005007	1.000000	-0.338481	-0.543351	-0.077221	-0.035322	0.081629	0.257307
Pclass	-0.035144	-0.338481	1.000000	0.131900	-0.369226	0.083081	0.018443	-0.549500


Jupyter Machine_learning_Assignment-4 Last Checkpoint: 12 minutes ago (autosaved) Python 3 (ipykernel)

File Edit View Insert Cell Kernel Widgets Help Trusted

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare
PassengerId	1.000000	-0.005007	-0.035144	0.042939	0.036847	-0.057527	-0.001652	0.012658
Survived	-0.005007	1.000000	-0.338481	-0.543351	-0.077221	-0.035322	0.081629	0.257307
Pclass	-0.035144	-0.338481	1.000000	0.131900	-0.369226	0.083081	0.018443	-0.549500
Sex	0.042939	-0.543351	0.131900	1.000000	0.093254	-0.114631	-0.245489	-0.182333
Age	0.036847	-0.077221	-0.369226	0.093254	1.000000	-0.308247	-0.189119	0.096067
SibSp	-0.057527	-0.035322	0.083081	-0.114631	-0.308247	1.000000	0.414838	0.159651
Parch	-0.001652	0.081629	0.018443	-0.245489	-0.189119	0.414838	1.000000	0.216225
Fare	0.012658	0.257307	-0.549500	-0.182333	0.096067	0.159651	0.216225	1.000000

```
In [53]: #Second form of visualizing correlation matrix using heatmap() from seaborn
sns.heatmap(matrix, annot=True, vmax=1, vmin=-1, center=0, cmap='vlag')
plt.show()
```

```
sns.heatmap(matrix, annot=True, vmax=1, vmin=-1, center=0, cmap='vlag')
plt.show()
```



In [54]: #Loaded data files test and train and merged files

```
train_raw = pd.read_csv('train.csv')
test_raw = pd.read_csv('test.csv')
```

```
Jupyter machine_learning_Assignment-4 Last Checkpoint: 13 minutes ago (autosaved) Logout
File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)
In [54]: #Loaded data files test and train and merged files
train_raw = pd.read_csv('train.csv')
test_raw = pd.read_csv('test.csv')
train_raw['train'] = 1
test_raw['train'] = 0
df = train_raw.append(test_raw, sort=False)
features = ['Age', 'Embarked', 'Fare', 'Parch', 'Pclass', 'Sex', 'SibSp']
target = 'Survived'
df = df[features + [target] + ['train']]
df['Sex'] = df['Sex'].replace(['female', 'male'], [0, 1])
df['Embarked'] = df['Embarked'].replace(['S', 'C', 'Q'], [1, 2, 3])
train = df.query('train == 1')
test = df.query('train == 0')

In [55]: # Drop missing values from the train set.
train.dropna(axis=0, inplace=True)
labels = train[target].values
train.drop(['train', target, 'Pclass'], axis=1, inplace=True)
test.drop(['train', target, 'Pclass'], axis=1, inplace=True)

In [56]: #Test and train split
X_train, X_val, Y_train, Y_val = train_test_split(train, labels, test_size=0.2, random_state=1)

In [57]: classifier = GaussianNB()
classifier.fit(X_train, Y_train)

Out[57]: GaussianNB
GaussianNB()
```

```
File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)
In [58]: y_pred = classifier.predict(X_val)

# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is', accuracy_score(Y_val, y_pred))

precision    recall  f1-score   support

      0.0      0.79      0.80      0.80         85
      1.0      0.70      0.69      0.70         58

 accuracy          0.76         143
macro avg          0.75          0.75         143
weighted avg          0.75          0.76         143

[[68 17]
 [18 40]]
accuracy is 0.7552447552447552

Question 3
(Glass Dataset)

1. Implement Naïve Bayes method using scikit-learn library.
a. Use the glass dataset available in Link also provided in your assignment.
b. Use train_test_split to create training and testing part.
2. Evaluate the model on testing part using score and classification_report(y_true, y_pred)

1. Implement linear SVM method using scikit library
a. Use the glass dataset available in Link also provided in your assignment.
```



2. Evaluate the model on testing part using score and

```
In [59]: glass=pd.read_csv("glass.csv")
glass.head()
```

```
Out[59]:
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
0	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0.0	0.0	1
1	1.51761	13.89	3.60	1.36	72.73	0.48	7.83	0.0	0.0	1
2	1.51618	13.53	3.55	1.54	72.99	0.39	7.78	0.0	0.0	1
3	1.51766	13.21	3.69	1.29	72.61	0.57	8.22	0.0	0.0	1
4	1.51742	13.27	3.62	1.24	73.08	0.55	8.07	0.0	0.0	1

```
In [60]: glass.corr().style.background_gradient(cmap="Reds")
```

```
Out[60]:
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
RI	1.000000	-0.191885	-0.122274	-0.407326	-0.542052	-0.289833	0.810403	-0.000386	0.143010	-0.164237
Na	-0.191885	1.000000	-0.273732	0.156794	-0.069809	-0.266087	-0.275442	0.326603	-0.241346	0.502898
Mg	-0.122274	-0.273732	1.000000	-0.481799	-0.165927	0.005396	-0.443750	-0.492262	0.083060	-0.744993
Al	-0.407326	0.156794	-0.481799	1.000000	-0.005524	0.325958	-0.259592	0.479404	-0.074402	0.598829
Si	-0.542052	-0.069809	-0.165927	-0.005524	1.000000	-0.193331	-0.208732	-0.102151	-0.094201	0.151565
K	-0.289833	-0.266087	0.005396	0.325958	-0.193331	1.000000	-0.317836	-0.042618	-0.007719	-0.010054
Ca	0.810403	-0.275442	-0.443750	-0.259592	-0.208732	-0.317836	1.000000	-0.112841	0.124968	0.000952
Ba	-0.000386	0.326603	-0.492262	0.479404	-0.102151	-0.042618	-0.112841	1.000000	-0.058692	0.575161
Fe	0.143010	-0.241346	0.083060	-0.074402	-0.094201	-0.007719	0.124968	-0.058692	1.000000	-0.188278
Type	-0.164237	0.502898	-0.744993	0.598829	0.151565	-0.010054	0.000952	0.575161	-0.188278	1.000000

machine_learning_Assignment-4 Last Checkpoint: 15 minutes ago (autosaved)

```
In [61]: sns.heatmap(matrix, annot=True, vmax=1, vmin=-1, center=0, cmap='vlag')
plt.show()
```

```
In [62]: features = ['RI', 'Na', 'Mg', 'Al', 'Si', 'K', 'Ca', 'Ba', 'Fe']
target = 'Type'
```



```
File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is', accuracy_score(y_val, y_pred))
```

	precision	recall	f1-score	support
1	1.00	0.95	0.97	19
2	0.57	1.00	0.73	12
3	0.00	0.00	0.00	6
5	0.00	0.00	0.00	1
6	0.00	0.00	0.00	1
7	0.75	0.75	0.75	4
accuracy			0.77	43
macro avg	0.39	0.45	0.41	43
weighted avg	0.67	0.77	0.70	43

```
[[18 1 0 0 0 0]
 [ 0 12 0 0 0 0]
 [ 0 5 0 0 0 1]
 [ 0 1 0 0 0 0]
 [ 0 1 0 0 0 0]
 [ 0 1 0 0 0 3]]
accuracy is 0.7674418604651163
```

Justification:
The Naive Bayes technique gave us a superior accuracy of 0.8372093023255814. While Linear SVM works better with linear regression logics, Naive Bayes analysis works better with probabilistic ideas. However in order to function more accurately, SVM needs a lot of data to train and test. Hence, compared to Linear SVM, Naive Bayes algorithm provides superior accuracy because of the volume of data.

The Naive Bayes technique gave us a superior accuracy of 0.8372093023255814. While Linear SVM works better with linear regression logics, Naive Bayes analysis works better with probabilistic ideas. However, in order to function more accurately, SVM needs a lot of data to train and test. Hence, compared to Linear SVM, Naive Bayes algorithm provides superior accuracy because of the volume of data.

Video link: <https://www.loom.com/share/081b23b321a54147a0fa5ae683848bb4>

Github link: <https://github.com/kadiresanjayreddy/machinelearningAssignment-4>